


AI/ML Performance Dashboard

AI/ML Dashboard Overview

This view combines model metrics, platform performance, trust indicators, and dataset readiness.



Main Parts Explained

1 Core ML

- Shows role, CIA, F1, and overall model accuracy for the main prediction services.
- Provides a fast health check for the primary ML models.

2 ML-based UABV

- Shows behavior-model accuracy, high-risk percentage, score difference, and contributing features.
- Helps assess the user-activity behavior model output quality.

3 RAG and LLM Performance

- Shows query volume, retrieval success, reasoning calls, and token consumption in grouped runtime panels.
- Helps compare retrieval quality and LLM usage efficiency together.

4 Human-in-the-Loop and Trust

- Shows manual corrections, override rates, and low-confidence predictions in one view.
- Highlights where human review and reliability controls affect the system.

5 Datasets / RAG / LLM

- Shows dataset volumes plus the active vector database, embedding model, and local LLM configuration.
- Provides readiness context for the data and runtime foundations behind the analytics stack.

This page is the operational source for AI/ML health monitoring, runtime visibility, and dataset readiness.